

Variational AutoEncoders

Latest Submission Grade 100%

1.

Question 1

For Variational AutoEncoders, which of the following are the correct operations performed in the *latent space*?

1 / 1 point

☐

encoder mean * encoder STDev * gaussian distribution

☒

encoder mean + encoder STDev * gaussian distribution

☐

encoder mean * encoder STDev + gaussian distribution

☐

encoder mean + encoder STDev + gaussian distribution

Correct

Correct!

2.

Question 2

Consider the following code, which is used in Variational AutoEncoder to represent the latent space. Fill in the missing piece of code.

(Note: Use shape as *shape=(batch, dim)*)

```
class Sampling(tf.keras.layers.Layer):  
    def call(self, inputs):  
        mu, sigma = inputs  
        batch = tf.shape(mu)[0]  
        dim = tf.shape(mu)[1]  
        epsilon = # YOUR CODE HERE  
        return mu + tf.exp(0.5 * sigma) * epsilon
```

1 / 1 point

tf.keras.backend.random_normal(shape=(batch, dim))

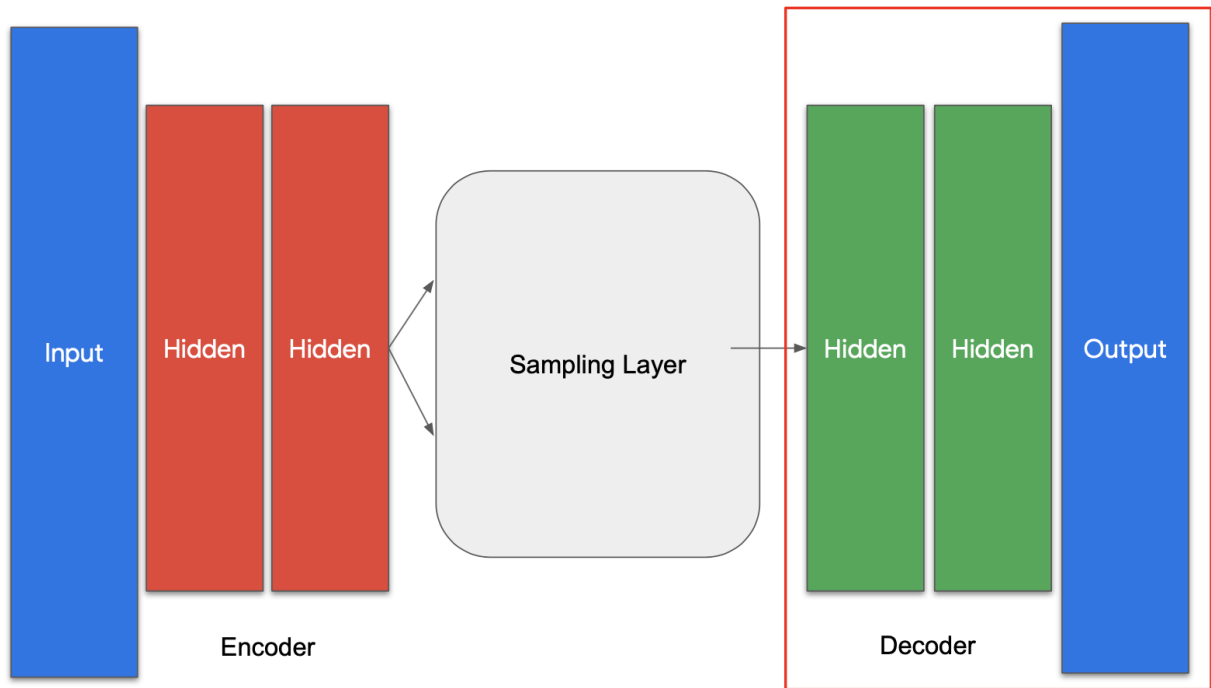
Correct

Correct!

3.

Question 3

When building the architecture for the decoder for a *convolutional Variational AutoEncoder*, what type of layers will you use ? Below is a screenshot of the code with # layer name # written in place of the actual layer that you would use. What goes in place of # layer name #?



```
def decoder_layers(inputs, conv_shape):
    units = conv_shape[1] * conv_shape[2] * conv_shape[3]
    x = tf.keras.layers.Dense(units, activation = 'relu',
                               name="decode_dense1")(inputs)
    x = tf.keras.layers.BatchNormalization()(x)

    x = tf.keras.layers.Reshape((conv_shape[1], conv_shape[2], conv_shape[3]),
                                name="decode_reshape")(x)
    x = tf.keras.layers.# layer name #(filters=64, kernel_size=3, strides=2,
                          padding='same', activation='relu',
                          name="decode_conv2d_2")(x)
    x = tf.keras.layers.BatchNormalization()(x)

    x = tf.keras.layers.# layer name #(filters=32, kernel_size=3, strides=2,
                          padding='same', activation='relu',
                          name="decode_conv2d3")(x)
    x = tf.keras.layers.BatchNormalization()(x)

    x = tf.keras.layers.# layer name #(filters=1, kernel_size=3, strides=1,
                          padding='same', activation='sigmoid', name="decode_final")(x)

    return x
```

1 / 1 point



Conv2DTranspose



Conv2D



MaxPooling2D.



Global AveragePooling2D

Correct

Correct! This will help you invert the convolutional filters applied during encoding.

4.

Question 4

Fill in the missing code for Kullback-Leibler cost function.

```
def kl_reconstruction_loss(inputs, outputs, mu, sigma):  
    kl_loss = # YOUR CODE HERE  
    return tf.reduce_mean(kl_loss) * - 0.5
```

1 / 1 point



`kl_loss = 1 + sigma - tf.square(mu) - tf.math.exp(sigma)`



`kl_loss = sigma - tf.square(mu) - tf.math.exp(sigma)`



`mu - tf.square(sigma) - tf.math.exp(mu)`



`kl_loss = 1 + mu - tf.square(sigma) - tf.math.exp(mu)`

Correct

Correct!